

Iterative Rigid Body Solvers

Richard Tonge

Principal Software Engineer, NVIDIA



Arena destruction demo





Demo credits

Matthias Müller-Fischer
Nuttapong Chentanez
Tae-Yong Kim
Aron Zoellner
Kevin Newkirk

Hawken Demo

SIGGRAPH paper

Mass Splitting for Jitter-Free Parallel Rigid Body Simulation

Richard Tonge, Feodor Benevolenski, Andrey Voroshilov

Mass Splitting for Jitter-Free Parallel Rigid Body Simulation

Richard Tonge
NVIDIA¹ Feodor Benevolenski
NVIDIA Andrey Voroshilov
NVIDIA



Figure 1: Left: Red cube with 4000 contacts coming to rest on a blue surface during mass splitting on an NVIDIA GT200 at over 600 FPS. Middle: Metal spinning used in a real-time physics simulation. The debris has irregular shapes and large mass ratios. Right: A complex scene with many objects interacting in real-time in a video game.

Abstract

Rigid body dynamics is widely used in applications ranging from real-time games to scientific simulations. Plans of objects are naturally common. Some of the most visually interesting simulations involve objects falling or crashing onto other objects. In general, these can generate large piles of debris. In mechanical engineering some of the most computationally challenging problems involve simulating the motion of such systems in real time on commodity hardware. This requires static collision detection, accurate dynamics, friction and reacting contact, which presents many challenges.

In large or real-time simulations, the computation cost can be prohibitive due to the number of contacts between bodies. In general we must use iterative methods, accumulating the iterations before convergence. We propose projected faculty by dividing each body mass into the effective mass by the number of contacts acting on the body. This allows us to parallelize the simulation and to estimate the method by solving contacts in blocks, providing quick convergence compared with PGS when solving using block projection. We present results for the authors' own implementation, which can simulate a pile of 5000 objects in real-time playing at over 600 FPS.

CR Categories: I.3.3 [Computer Graphics]: Computational Geometry and Object Modeling; Physics-based modeling; I.3.5 [imulation and Modeling]; I.4.0 [Immersive—Interactive]

Keywords: rigid bodies, non-smooth dynamics, contact, division

Links: [DOI](#) [PDF](#)

¹ Email: rtonge@nvidia.com

Due to the widespread availability of multi-core CPUs and GPUs, parallel rigid body simulation has become increasingly popular. PGS has limited parallelism, as updates to bodies having multiple contacts must be serialized to ensure they are not lost and to prevent numerical instability due to the accumulation of errors. This serialization changes the order in which constraints are processed, often changing it dramatically from frame to frame. Typically, the order of constraint processing is determined by the order of the operations. For these reasons it is very challenging to avoid pairing with parallel PGS, and due to the serialization it performance drops two-fold as more threads are added.

In real-time applications such as games, the designer does not know in advance what the player is going to do and which objects are

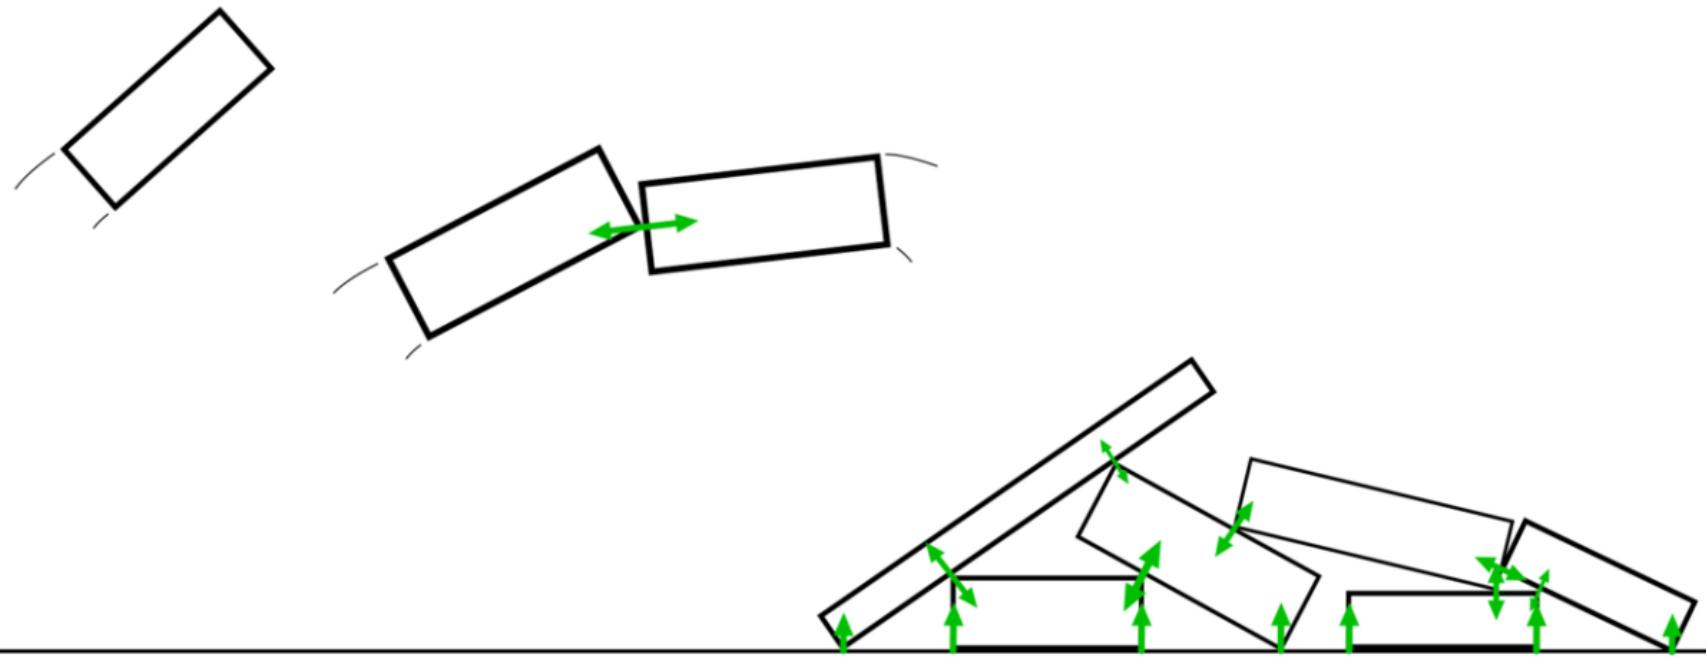
Contents

- What do we need rigid bodies to do?
- Single core solver
- Parallel solver
- Results

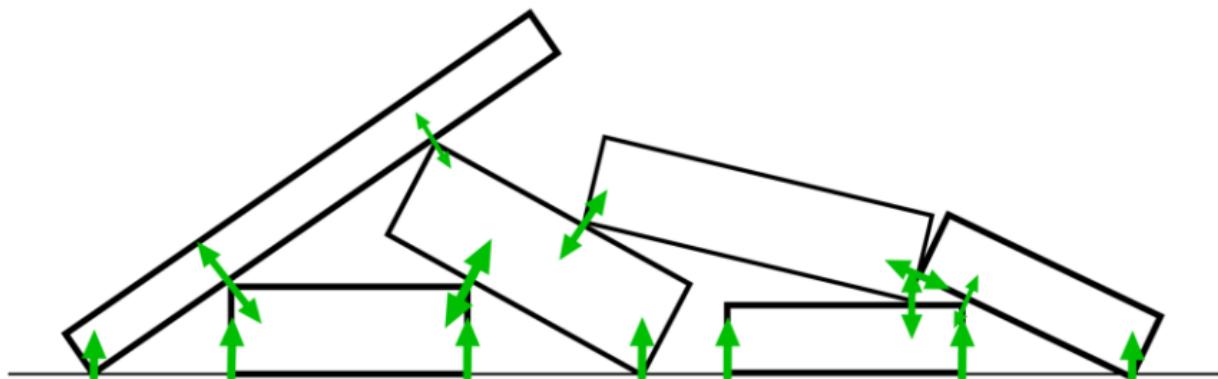
Section 1

What are rigid bodies supposed to do?





Stable Piles = Balanced forces



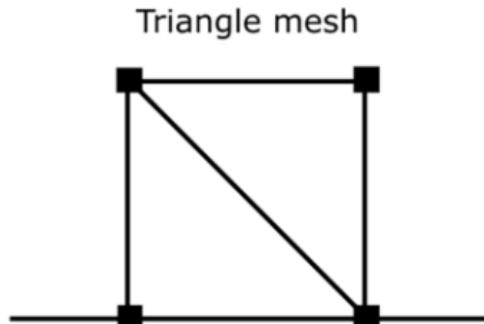
Contacts affect each other



Section 2

How to make rigid bodies do the right thing
using a solver (single core)

Rigid Body Basics



Rigid body coordinates of
triangle mesh



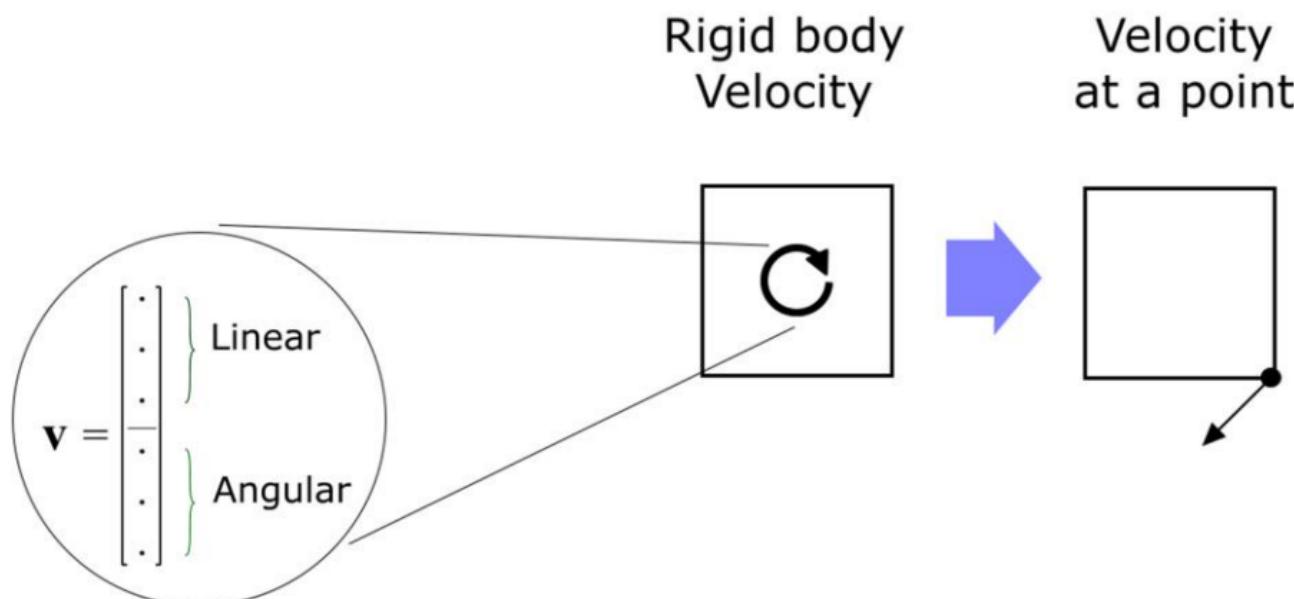
$$\mathbf{x} = \begin{bmatrix} \cdot \\ \cdot \\ \vdots \\ \cdot \\ \cdot \\ \cdot \end{bmatrix}$$

Position

Orientation

The equation $\mathbf{x} = \begin{bmatrix} \cdot \\ \cdot \\ \vdots \\ \cdot \\ \cdot \\ \cdot \end{bmatrix}$ represents the state of a rigid body. The first three rows represent the Position (center of mass), and the last three rows represent the Orientation (represented by a 3x3 rotation matrix).

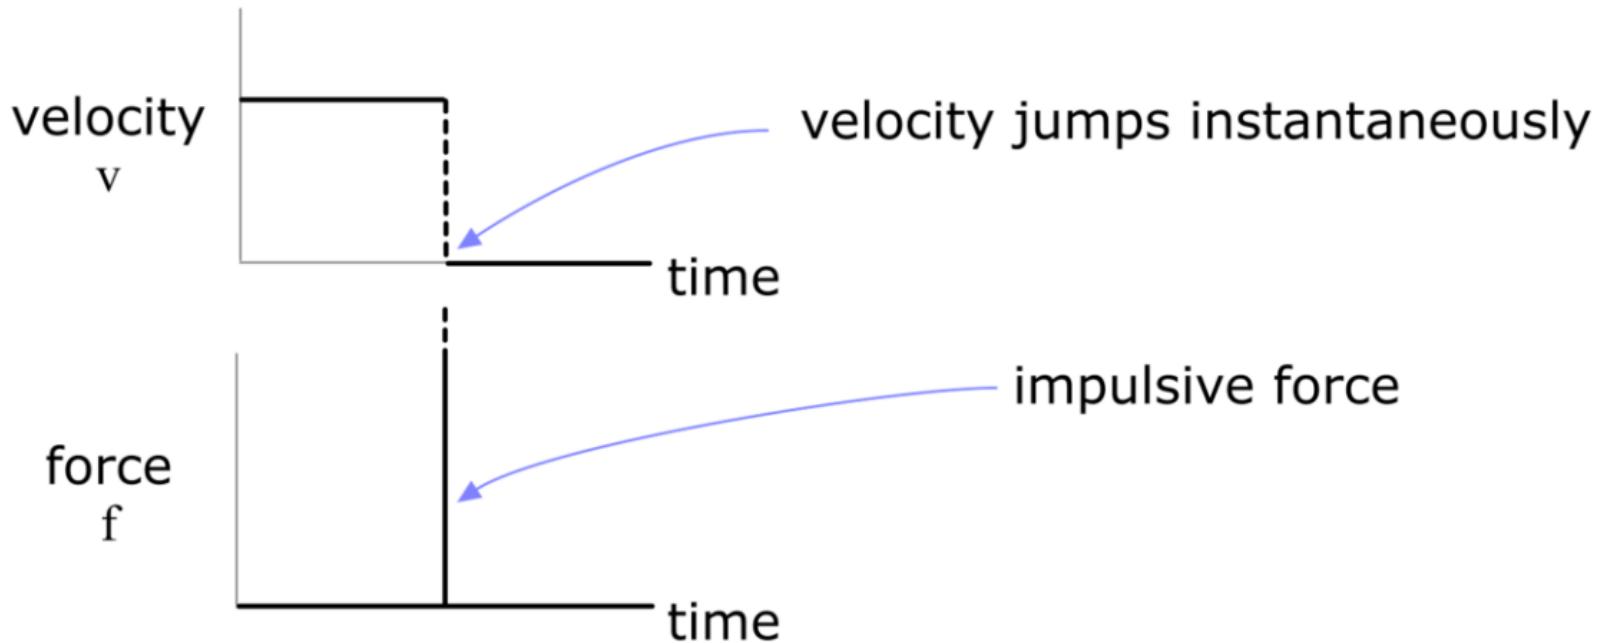
Velocities



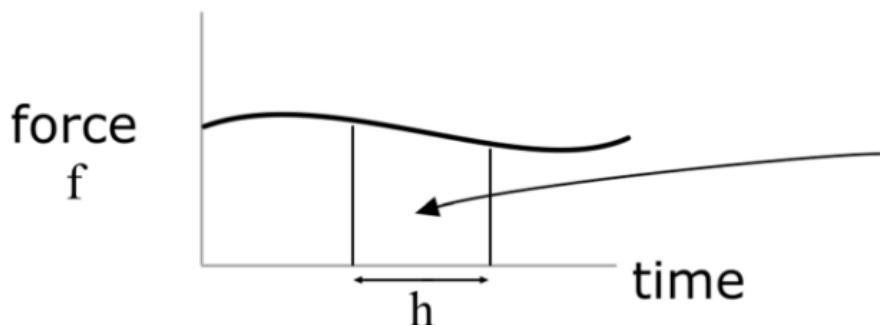
Use impulses and velocities



What is an impulse?



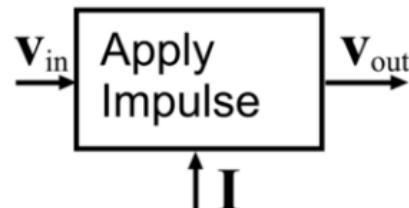
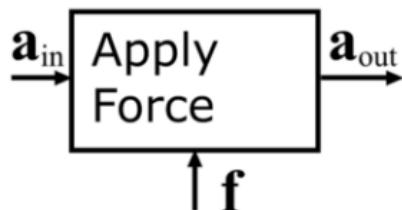
What is an impulse?



impulse is area under
force-time graph

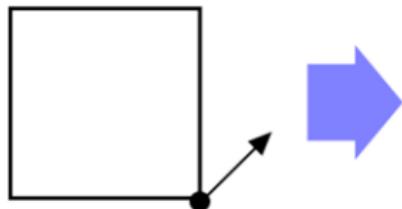
(for constant force, $I = hf$)

Impulse application

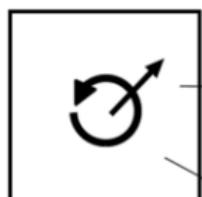


Impulses

Impulse
applied at point



Rigid Body
Impulse



$$\lambda = \begin{bmatrix} \cdot \\ \cdot \\ \cdot \\ \cdot \\ \cdot \end{bmatrix}$$

The equation shows a column vector λ with five entries, each represented by a dot. To the right of the vector, a green brace groups the first four entries, labeled "Linear". Below the vector, another green brace groups the last two entries, labeled "Angular".

Making a rigid body solver

Contacts

From Collision
Detection

Rigid Body
Coordinates



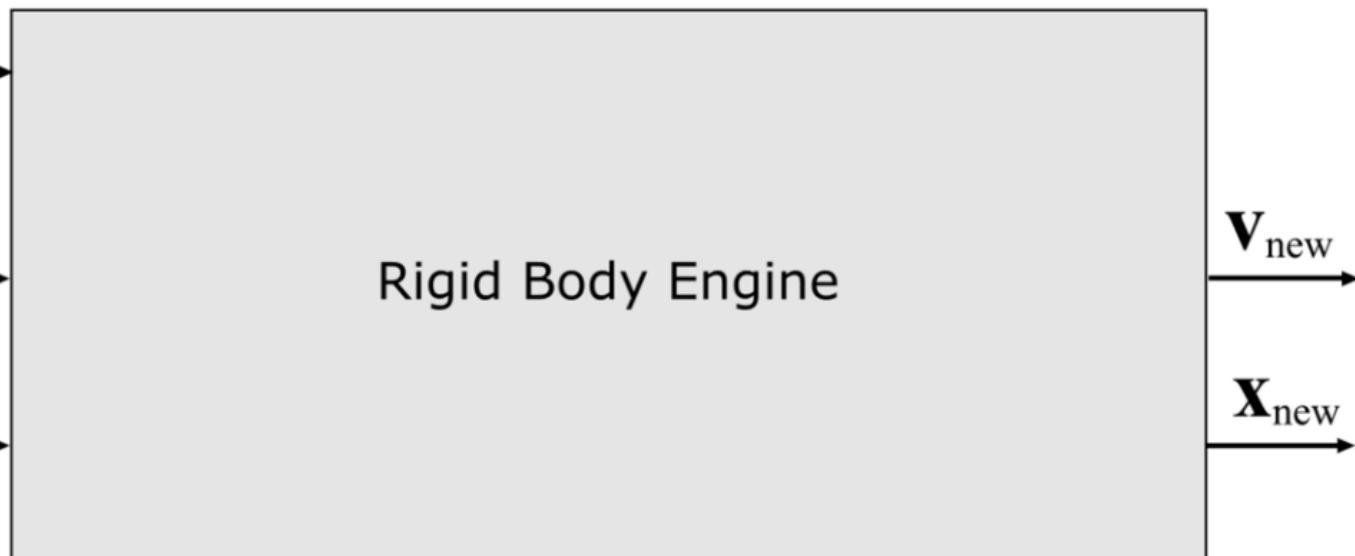
V_{old}

X_{old}

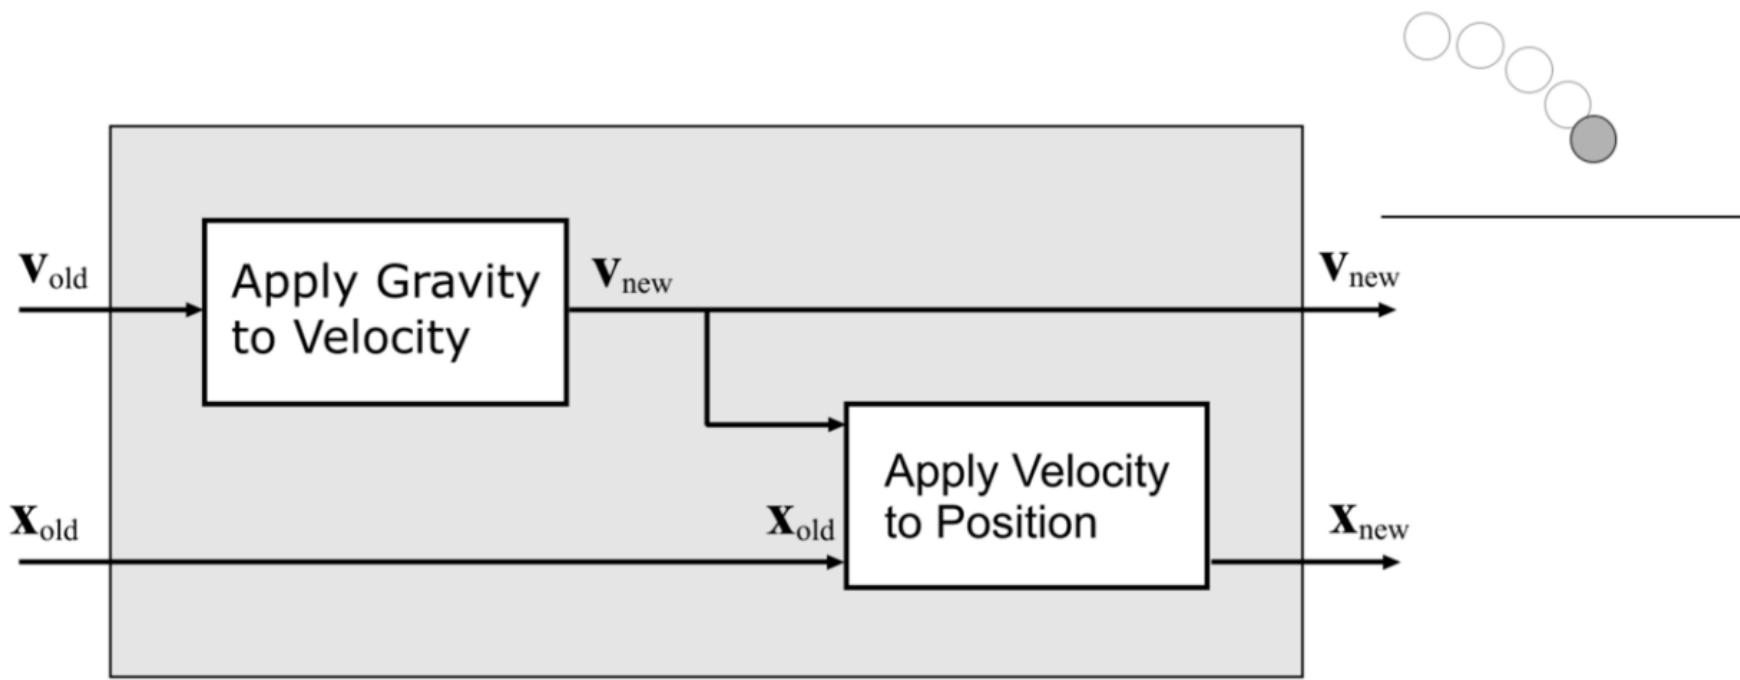
Rigid Body Engine

V_{new}

X_{new}



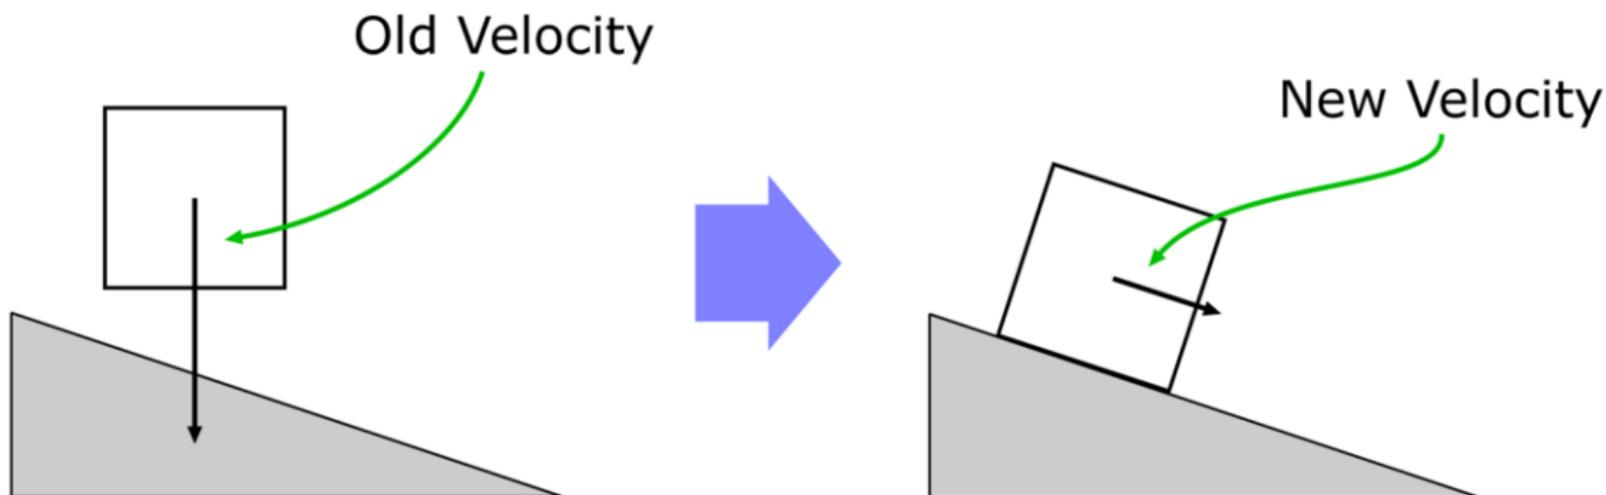
Moving a Body Without Collisions



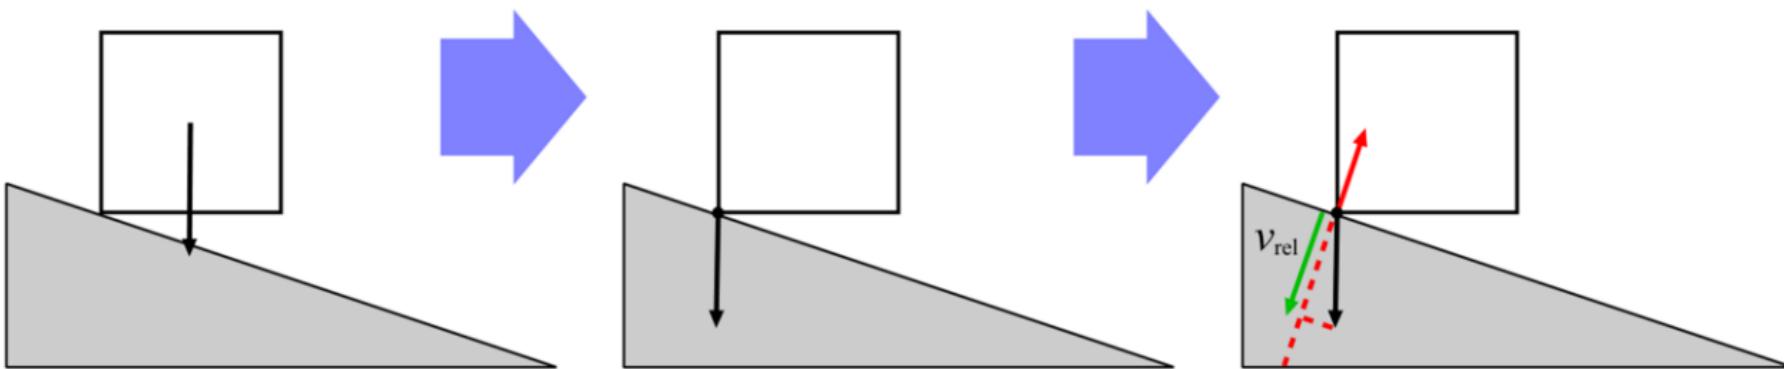
Adding a single contact



Contact at the velocity level



Velocity at the contact

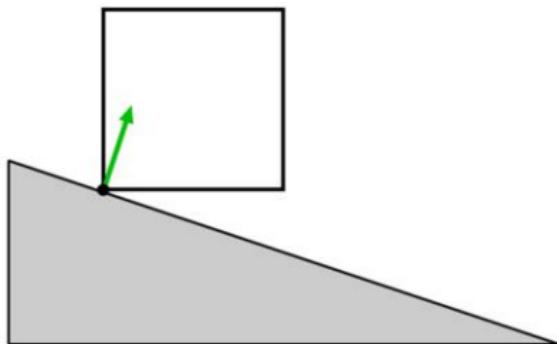


Calculating the impulse

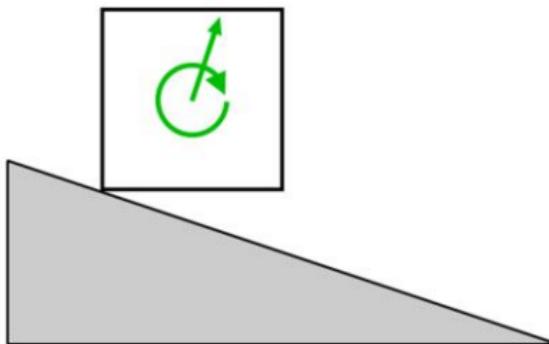


Converting impulse to RB coords

Impulse
At Point

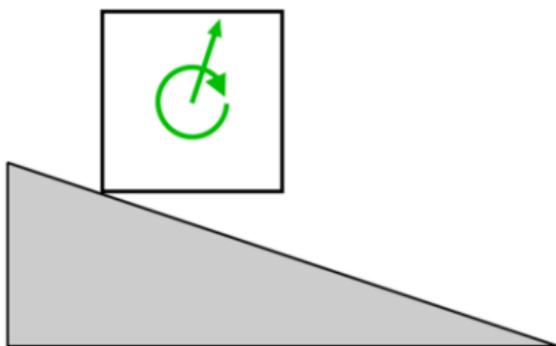


Rigid Body
Impulse

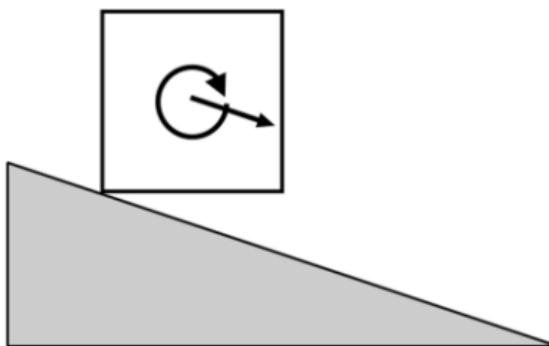


Applying the impulse

Rigid Body
Impulse

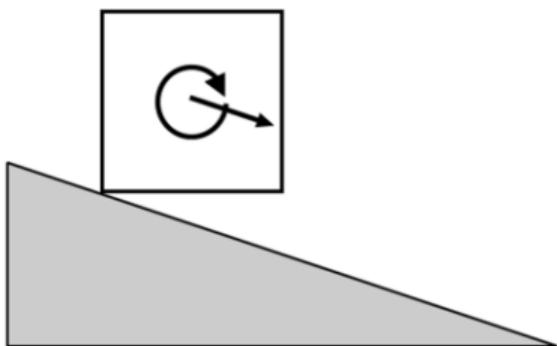


Rigid Body
Velocity

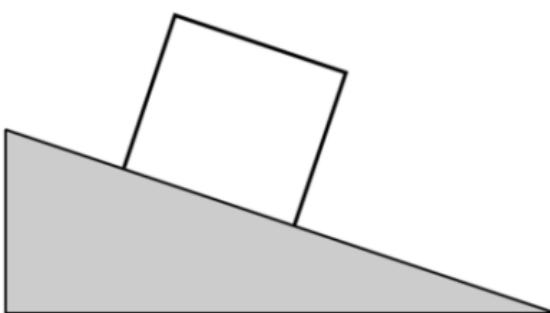


Applying the velocity

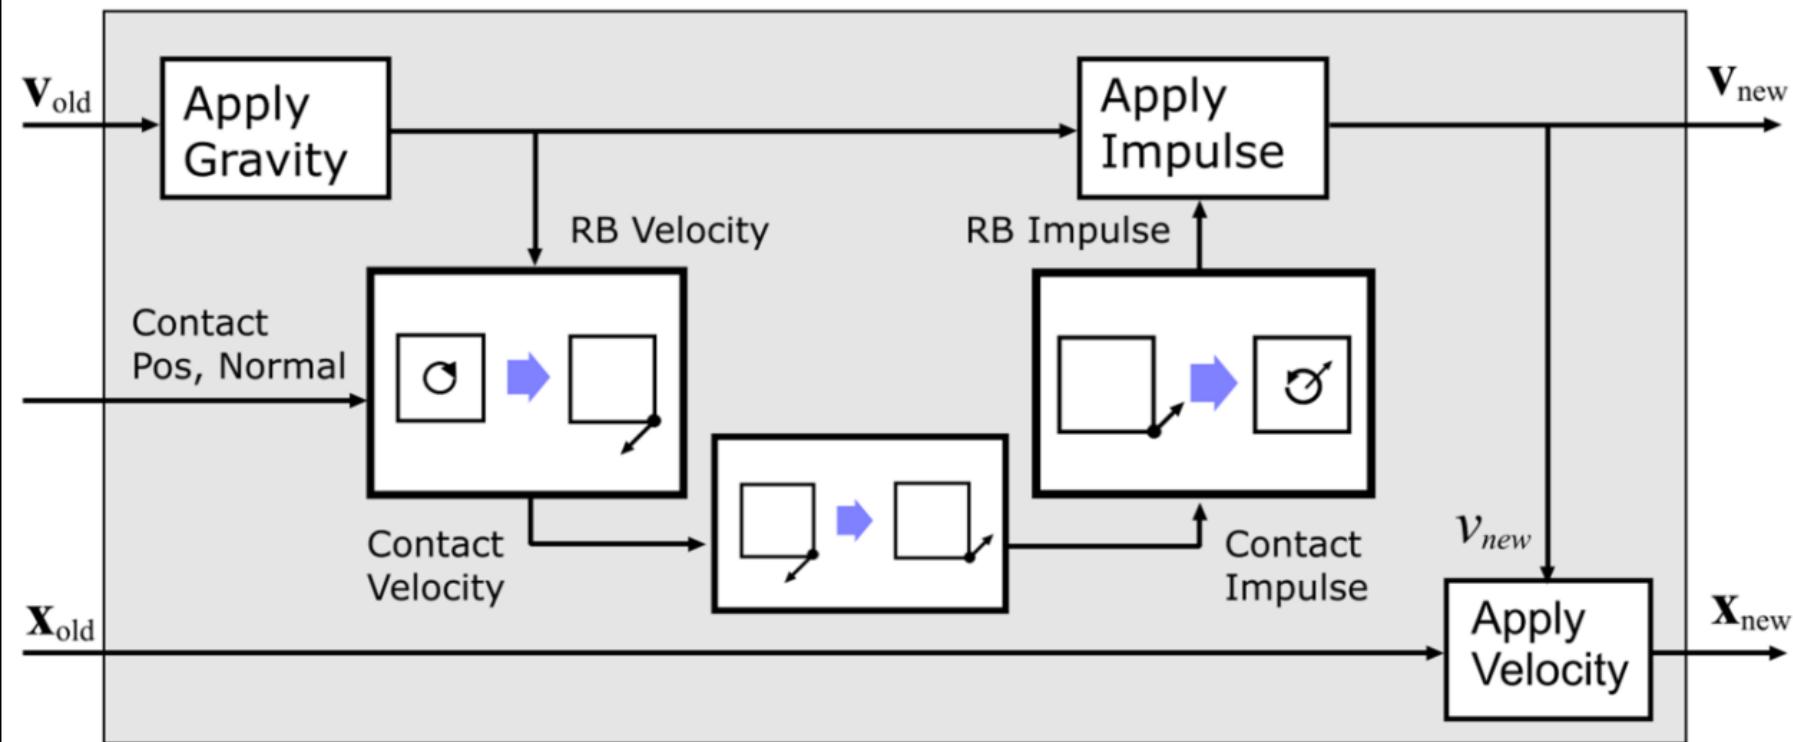
Rigid Body
Velocity



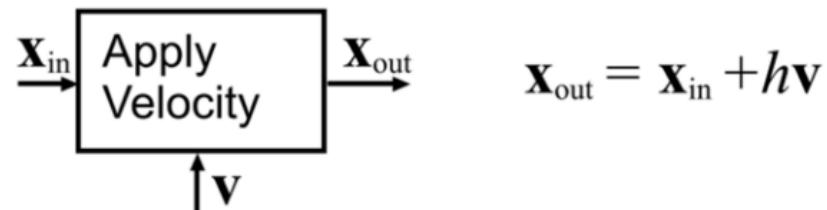
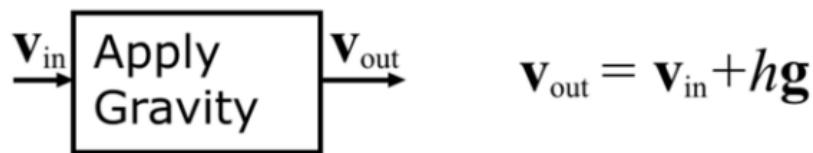
Updated
Pose



Putting it all together

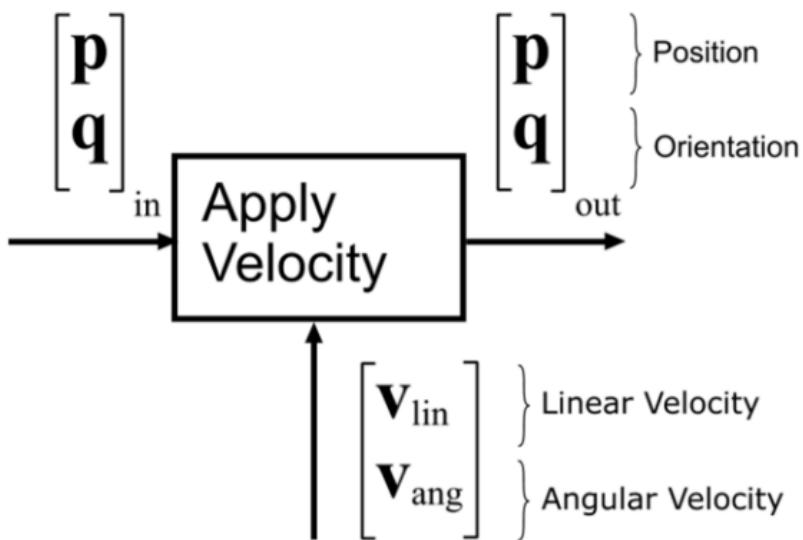


Summary of Operations 1



h is the timestep size, 1/60 seconds for 60Hz

Rotation Complicates Things



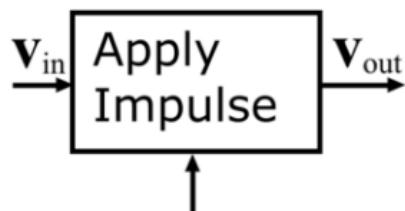
$$\mathbf{p}_{out} = \mathbf{p}_{in} + h\mathbf{v}_{lin}$$

$$\mathbf{q}_{out} = [\cos(\theta/2), \mathbf{a} \sin(\theta/2)]\mathbf{q}_{in}$$

$$\mathbf{a} = \mathbf{v}_{ang}/|\mathbf{v}_{ang}|$$

$$\theta = |h\mathbf{v}_{ang}|$$

Apply Impulse



$$\mathbf{v}_{out} = \mathbf{v}_{in} + \mathbf{M}^{-1} \mathbf{I}$$

Rigid Body Impulse, \mathbf{I}

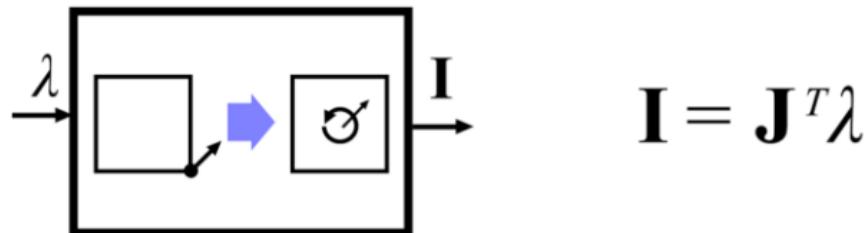
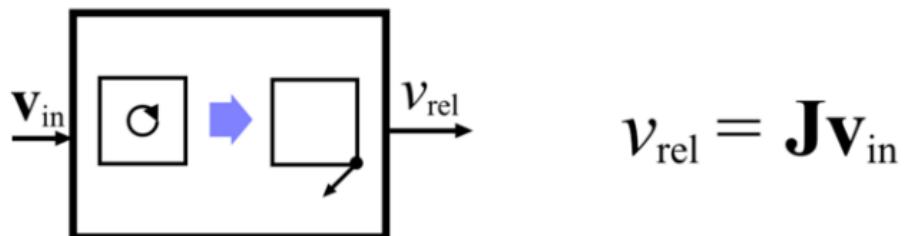
Mass Matrix

$$\mathbf{M} = \begin{bmatrix} m & & & \\ & m & & \\ & & m & \\ & \vdots & \vdots & \ddots \end{bmatrix}$$

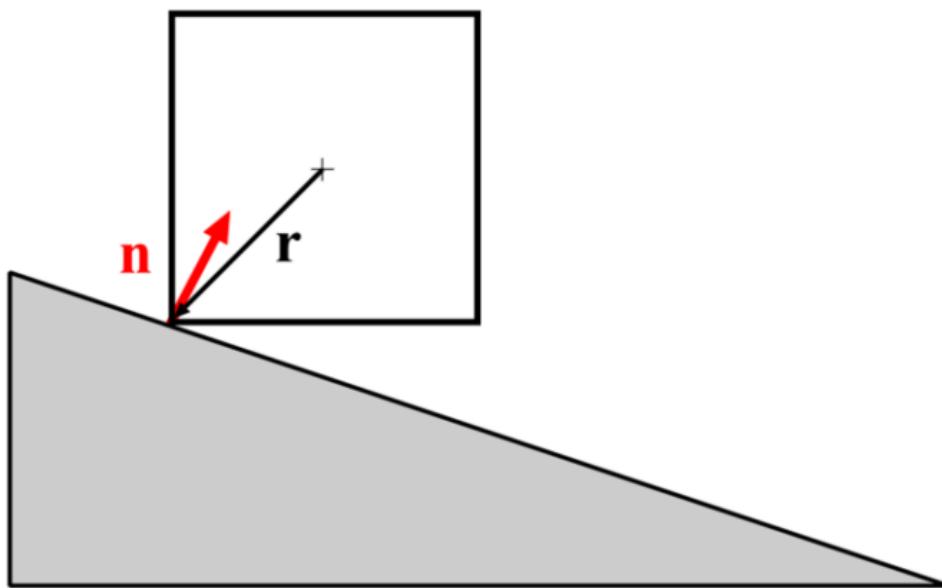
Inertia Tensor

The diagram shows a 4x4 matrix with 'm' in the diagonal entries. The bottom-right 2x2 submatrix is circled in green and labeled "Inertia Tensor".

Summary of operations 2

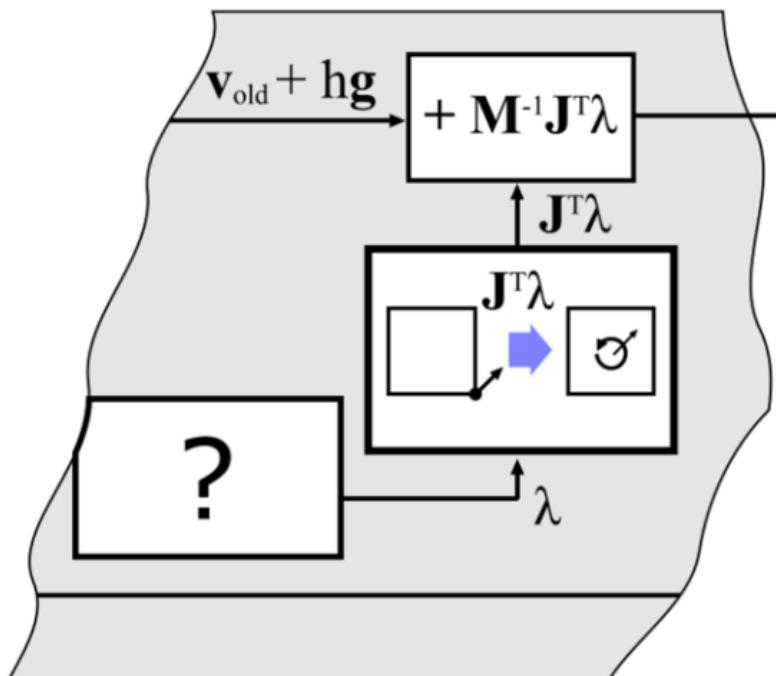


What is \mathbf{J} ?



$$\mathbf{J} = [\mathbf{n} \mid \mathbf{r} \times \mathbf{n}]$$

Calculating the impulse (λ)



Final velocity in terms of λ

$$\mathbf{v}_{\text{new}} = \mathbf{v}_{\text{old}} + \mathbf{hg} + \mathbf{M}^{-1} \mathbf{J}^T \lambda$$



Final relative velocity

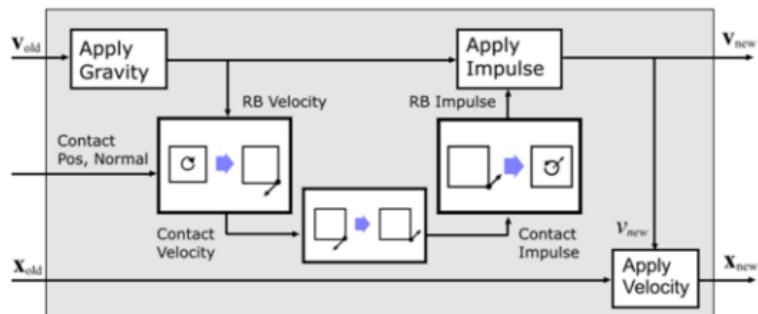
$$\mathbf{J}\mathbf{v}_{\text{new}} = \mathbf{J}(\mathbf{v}_{\text{old}} + \mathbf{hg}) + \mathbf{J}\mathbf{M}^{-1} \mathbf{J}^T \lambda$$



Final relative velocity is zero when:

$$\lambda = -(\mathbf{J}\mathbf{M}^{-1} \mathbf{J}^T)^{-1} \mathbf{J}(\mathbf{v}_{\text{old}} + \mathbf{hg})$$

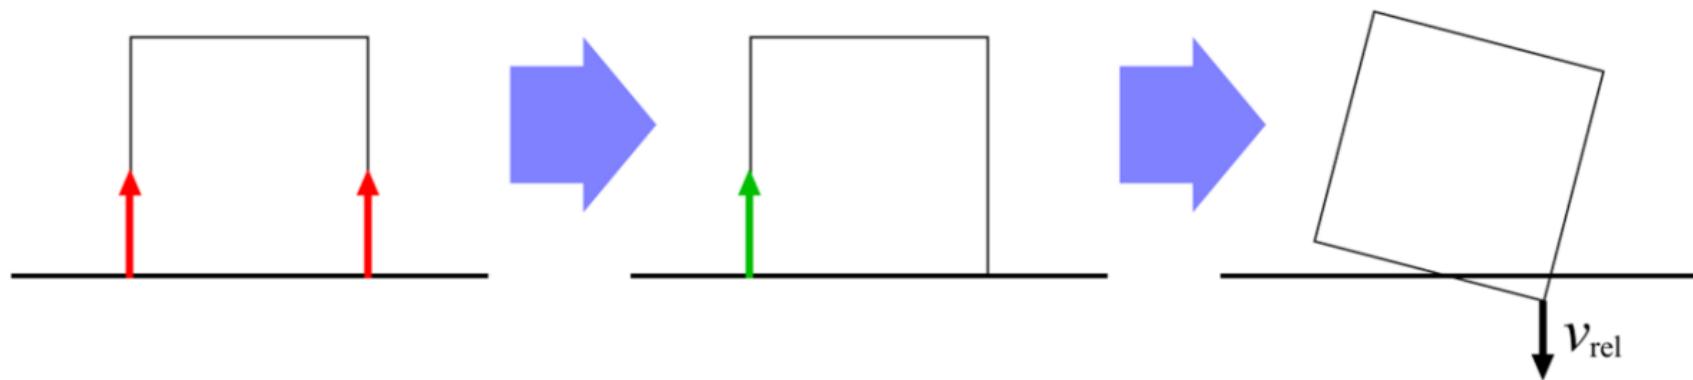
Single Contact Algorithm



=

```
solveSingle(x, v, J, M, h, g)
{
    λ = (-(JM⁻¹JT)⁻¹J(v + hg))⁺
    v = v + M⁻¹JTλ
    x = x + hv
}
```

Multiple Contacts



Two Contacts

Single Impulse

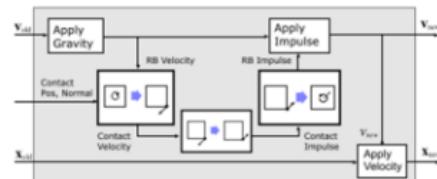
Other Contact
Worse

Multiple contacts



$$0 \leq \mathbf{v}_{\text{rel}} \perp 0 \leq \lambda$$

Model first



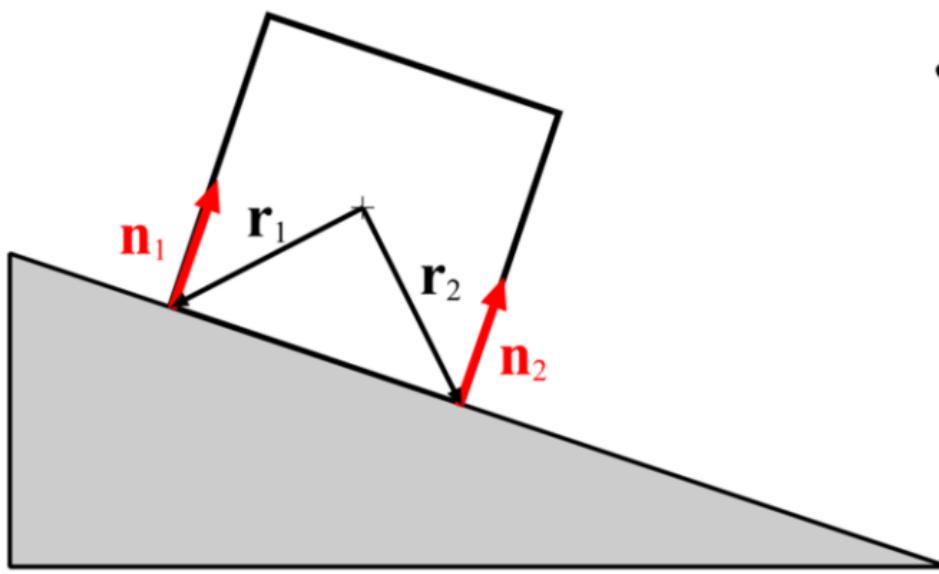
Algorithm second

The value of knowing the model

- Something to test against
- Convergence
- Peace of mind when debugging
- Other solutions for same model



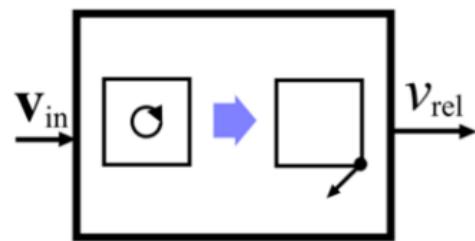
What is \mathbf{J} ?



$$\begin{aligned}\mathbf{J} &= \left[\begin{array}{c|c} \mathbf{n}_1 & \mathbf{r}_1 \times \mathbf{n}_1 \\ \hline \mathbf{n}_2 & \mathbf{r}_2 \times \mathbf{n}_2 \end{array} \right] \\ &= \left[\begin{array}{ccc|ccc} \cdot & \cdot & \cdot & \cdot & \cdot & \cdot \\ \hline \cdot & \cdot & \cdot & \cdot & \cdot & \cdot \end{array} \right]\end{aligned}$$



Operations still work with new \mathbf{J}



$$\mathbf{v}_{\text{rel}} = \mathbf{J}\mathbf{v}_{\text{in}}$$

$$\begin{aligned} v_{\text{rel}} \text{ at contact 1 } [\cdot] &= \left[\begin{array}{cccc|cc} \cdot & \cdot & \cdot & \cdot & \cdot & \cdot \\ \hline \cdot & \cdot & \cdot & \cdot & \cdot & \cdot \end{array} \right] \begin{bmatrix} \cdot \\ \vdots \\ \cdot \\ \vdots \\ \cdot \end{bmatrix} \\ v_{\text{rel}} \text{ at contact 2 } [\cdot] &= \left[\begin{array}{cccc|cc} \cdot & \cdot & \cdot & \cdot & \cdot & \cdot \\ \hline \cdot & \cdot & \cdot & \cdot & \cdot & \cdot \end{array} \right] \begin{bmatrix} \cdot \\ \vdots \\ \cdot \\ \vdots \\ \cdot \end{bmatrix} \end{aligned}$$



System is a matrix equation?



No



$$\lambda = \text{LCP}(A, b)$$

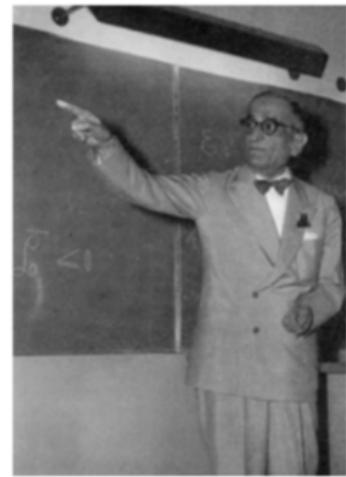
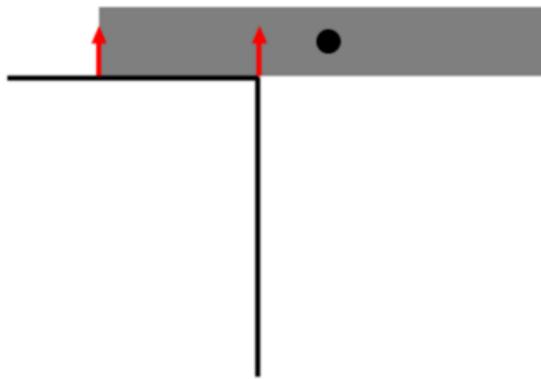
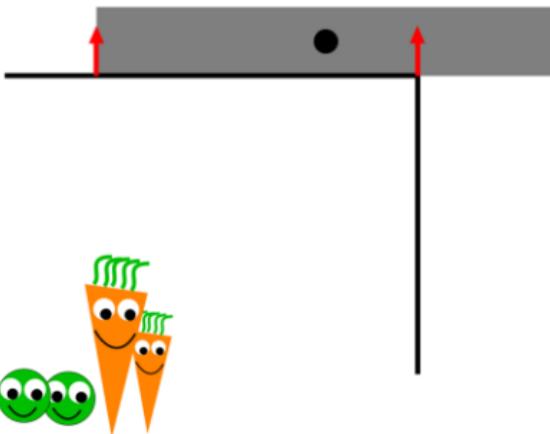


Why?

Contacts can break



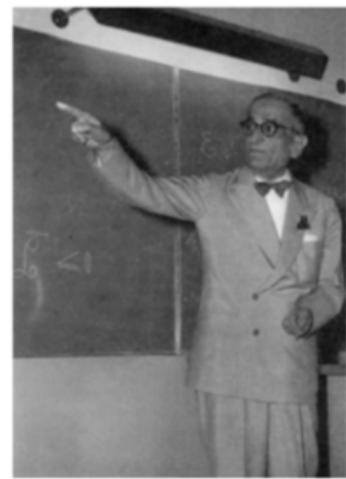
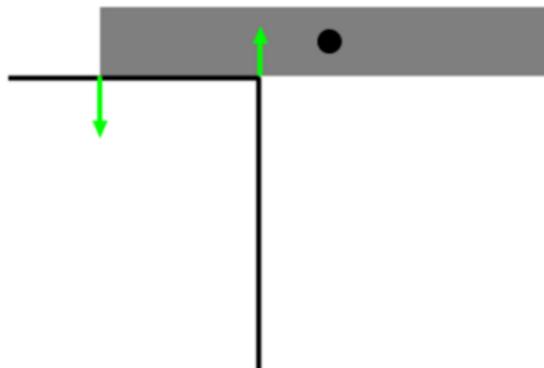
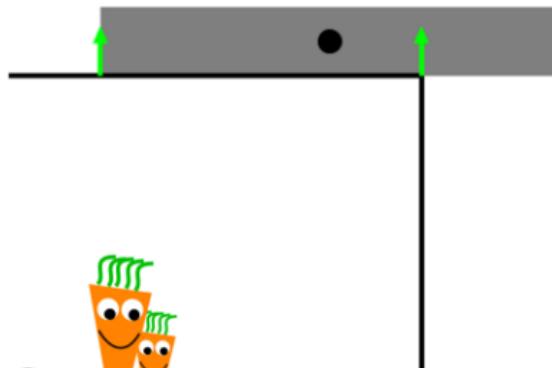
When Should Contacts Break?



Antonio Signorini



Linear system would eliminate velocity at all contacts



Antonio Signorini



Bar won't fall

The Signorini Conditions:

$0 \leq v_{\text{rel}}$ Every relative velocity should
be zero or separating

$0 \leq \lambda$ Every contact impulse should
be non-attractive



Antonio Signorini

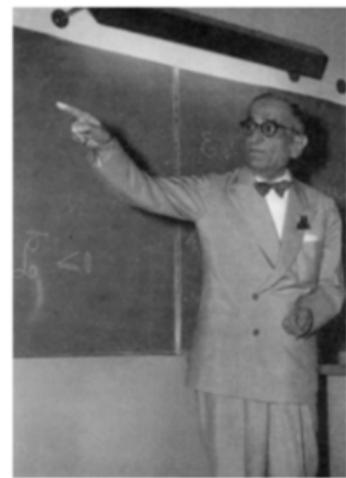


$$(v_{\text{rel}})_i = 0 \text{ or } \lambda_i = 0$$

No impulse at separating contacts

A compact way to write the three conditions
in one line of math:

$$0 \leq v_{\text{rel}} \perp 0 \leq \lambda$$



Antonio Signorini



The Final Model

$$\mathbf{M}\ddot{\mathbf{x}} = \mathbf{J}^T \boldsymbol{\lambda} + \mathbf{f}_e$$

$$\dot{\mathbf{x}} = \mathbf{v}$$

$$\mathbf{0} \leq \boldsymbol{\lambda} \perp \mathbf{0} \leq \mathbf{Jv}$$



Discretized model

$$\mathbf{M}(\mathbf{v}_{\text{new}} - \mathbf{v}_{\text{old}}) = \mathbf{J}^T \mathbf{z} + h\mathbf{f}_e$$

$$\mathbf{x}_{\text{new}} - \mathbf{x}_{\text{old}} = h\mathbf{v}_{\text{new}}$$

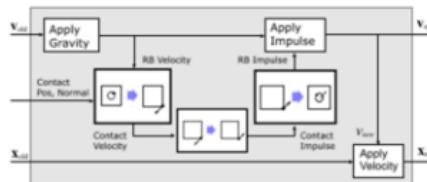
$$\mathbf{z} \geq 0 \perp \mathbf{J}\mathbf{v}_{\text{new}} \geq 0$$

$$\lambda = LCP\left(\mathbf{J}\mathbf{M}^{-1}\mathbf{J}^T, \mathbf{J}(\mathbf{v}_{\text{old}} + h\mathbf{g})\right)$$

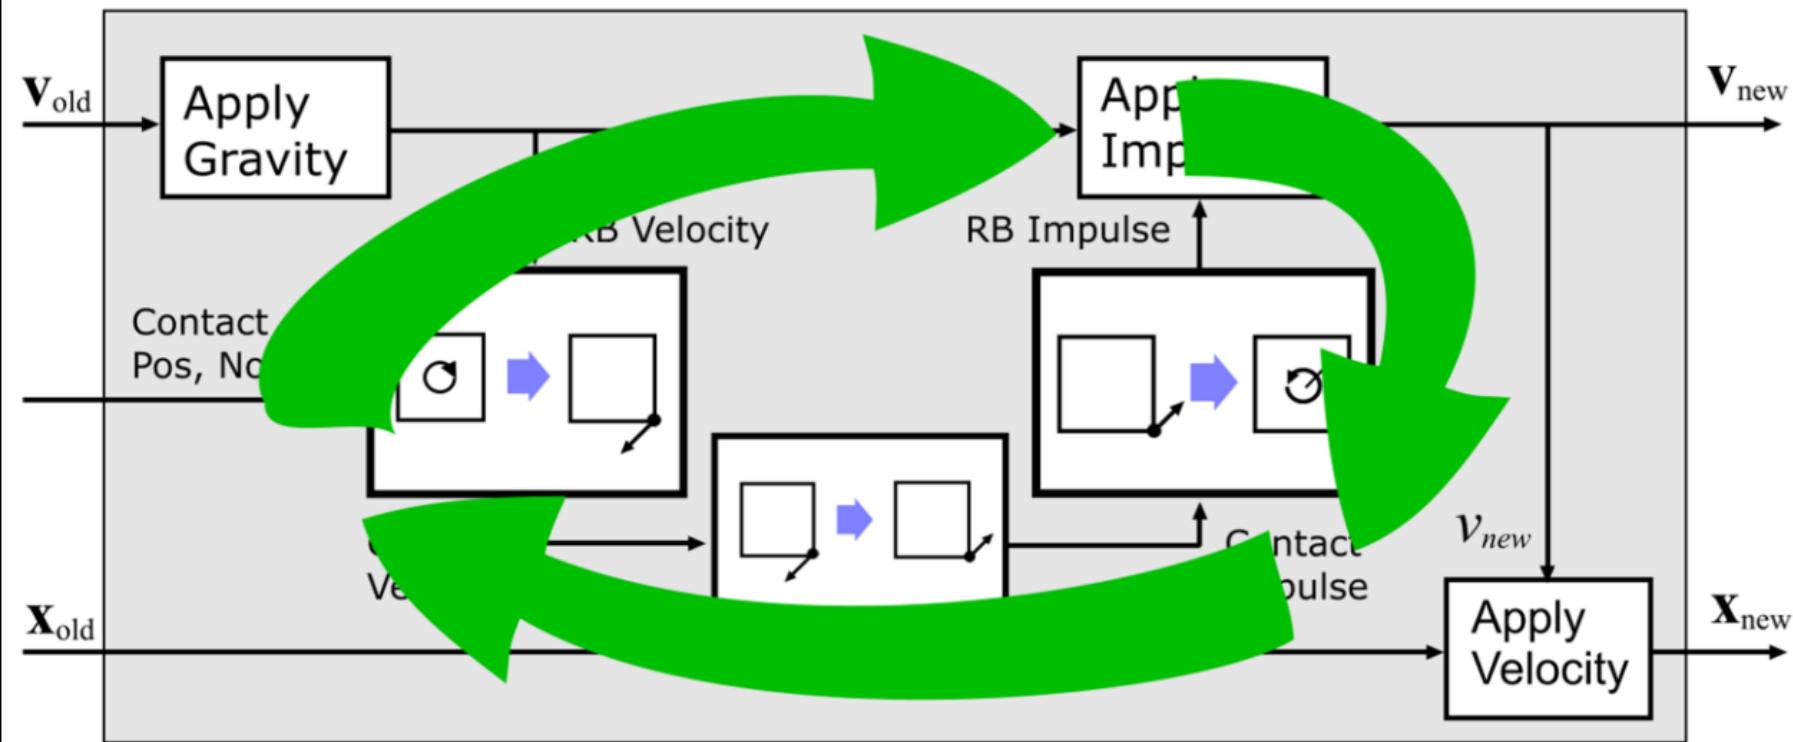
$$\mathbf{v}_{\text{new}} = \mathbf{v}_{\text{old}} + \mathbf{M}^{-1}\mathbf{J}^T \mathbf{z} + h\mathbf{g}$$

$$\mathbf{x}_{\text{new}} = \mathbf{x}_{\text{old}} + h\mathbf{v}_{\text{new}}$$

Just Sequentially Apply Impulses



Sequentially Applying Impulses



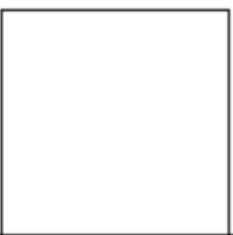
Initial idea to enforce Signorini

At each iteration,
if impulse (λ) is negative, set it to zero

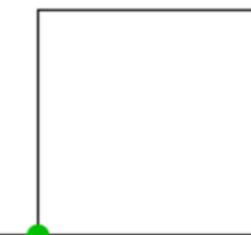
Potential problem



Initial impulse
too high



Try to correct...



...but -ve impulse
gets clamped

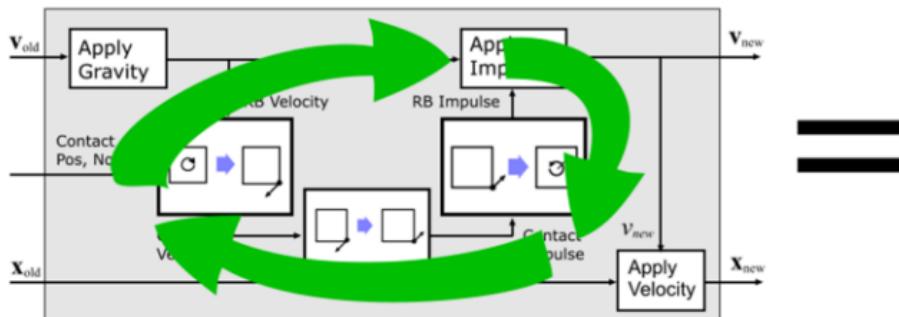


Impulse is always
too large

Solution: Clamp total impulse

- Keep impulse accumulator for each contact
- Clamp the accumulator at iteration, not the added impulse

Multiple Contact Algorithm



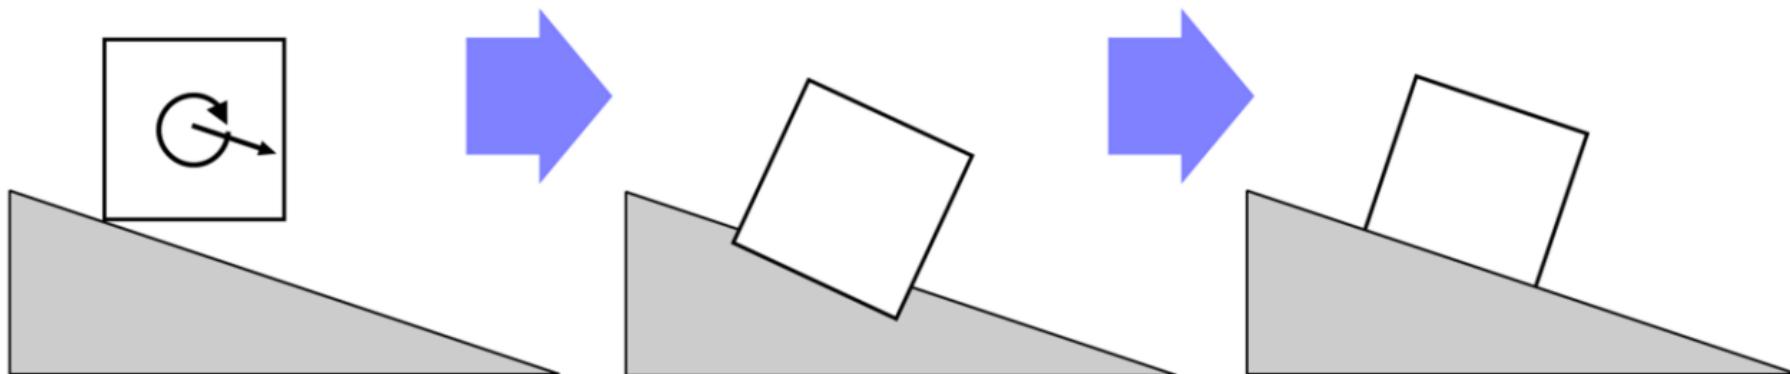
```
solveMultiple(x, v, J, M, h, g)
{
    for j=0; j<iterCnt; j++
    {
        for i=0; i<contactCnt; i++
        {
            t = λi
            m = JiM-1JiT
            λi = λi - (1/m)(Jv + b)
            λi = max(0, λi)
            v = v + M-1JT(λi - t)
        }
    }
    x = x + hv
```

Position projection

Rigid Body
Velocity

Fixed timestep,
moved into penetration

After position
projection



Position Projection

Let ϕ be the penetration.

Instead of requiring: $\mathbf{Jv} \geq \mathbf{0}$

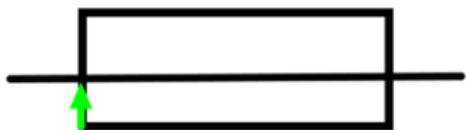
Require: $\mathbf{Jv} + (\gamma/h)\phi \geq \mathbf{0}$

In previous algorithm, set $\mathbf{b} = (\gamma/h) \phi$
(I use $\gamma = 0.8$)

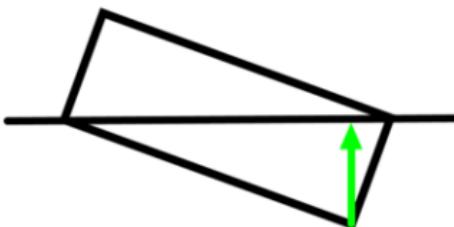
Section 3

How to make rigid bodies do the right thing
using a solver (multicore)

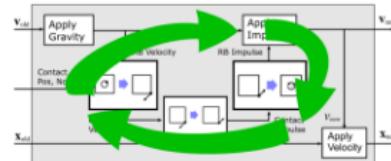
Existing solver method 1:



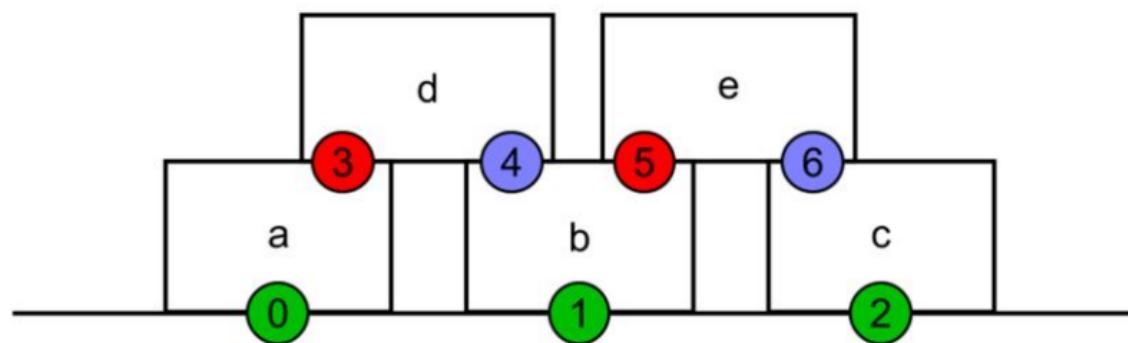
Penetrating configuration



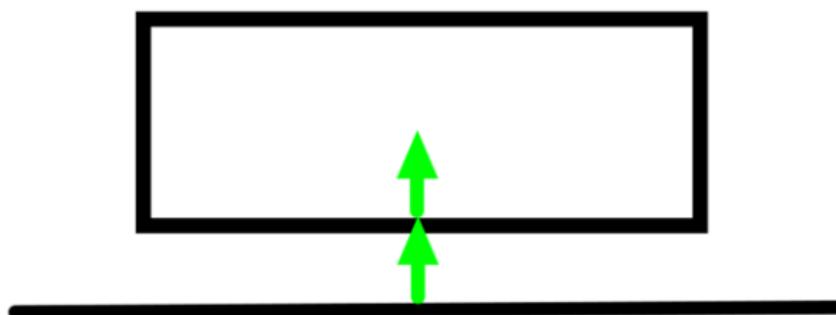
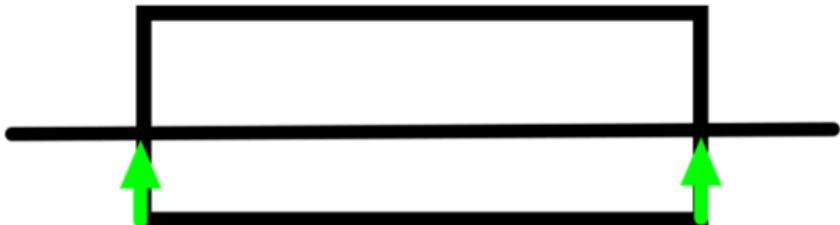
Rendered Frame



PGS coloring



Existing solver method 2:



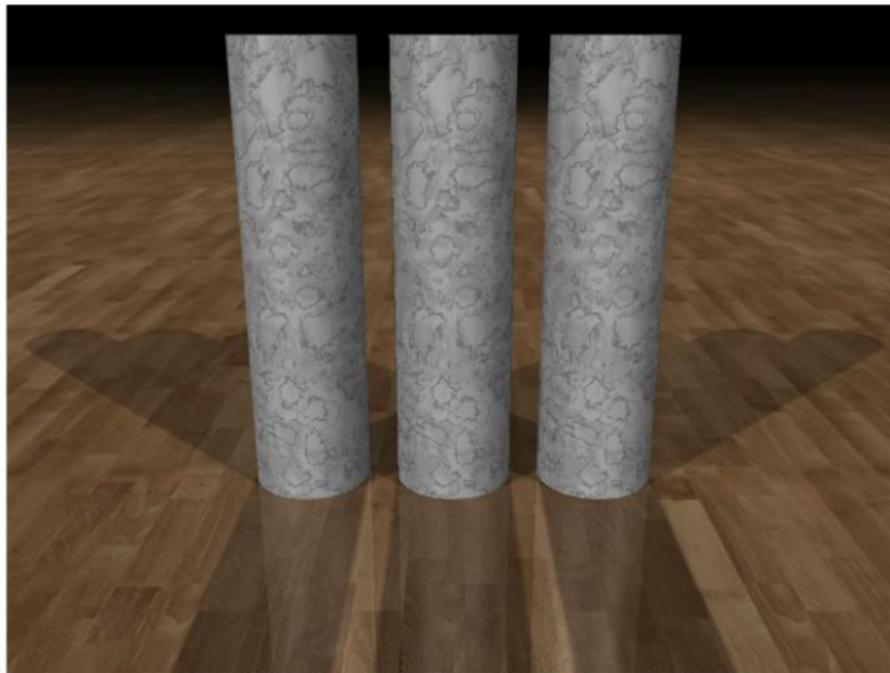
- Method 1 (Projected Gauss Seidel, PGS)

- Provably convergent ✓
- Limited parallelism ✗
- Jitters ✗
- Widely used

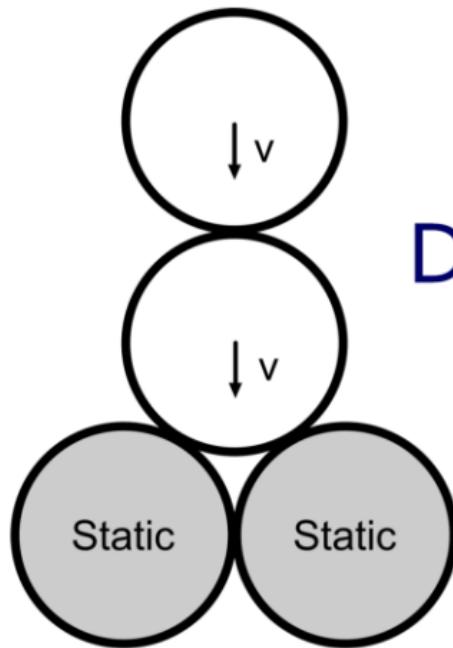
- Method 2 (Projected Jacobi)

- Maximally parallel ✓
- Jitter free ✓
- Non convergent in many cases ✗
- Converges slowly ✗
- Unusable in games

Parallel PGS Jitter

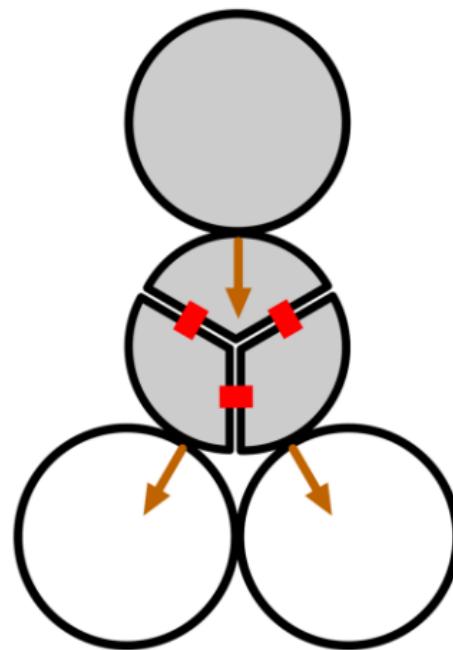


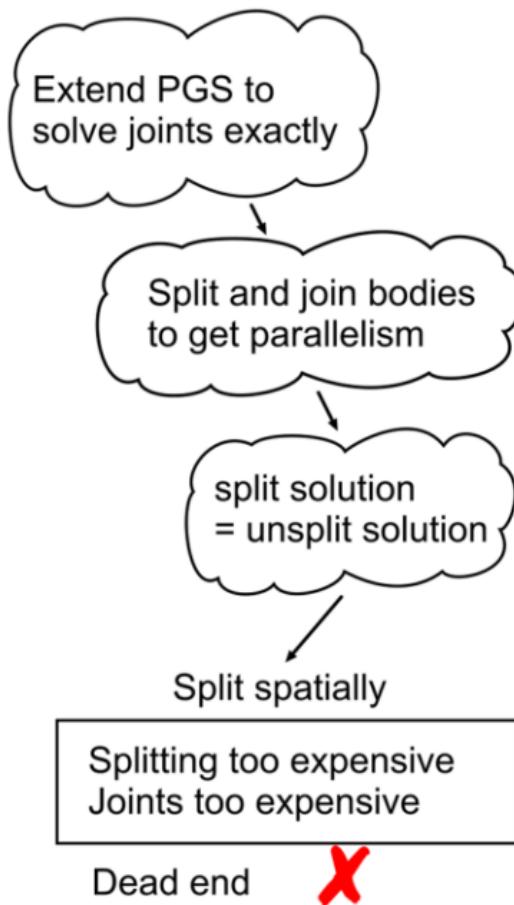
Example



Doesn't converge with Jacobi

First idea: Spatial splitting

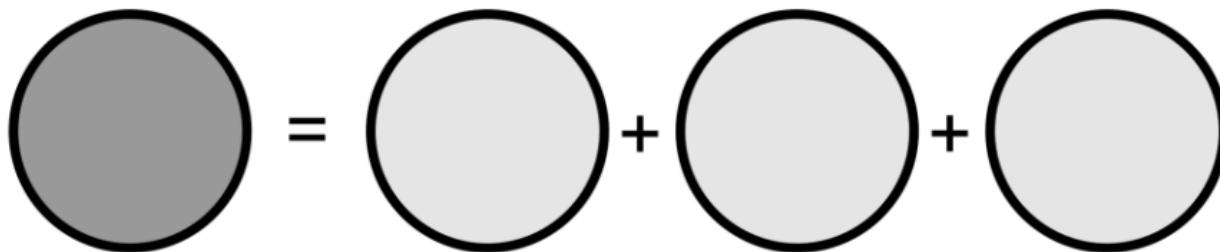




New idea: Mass splitting

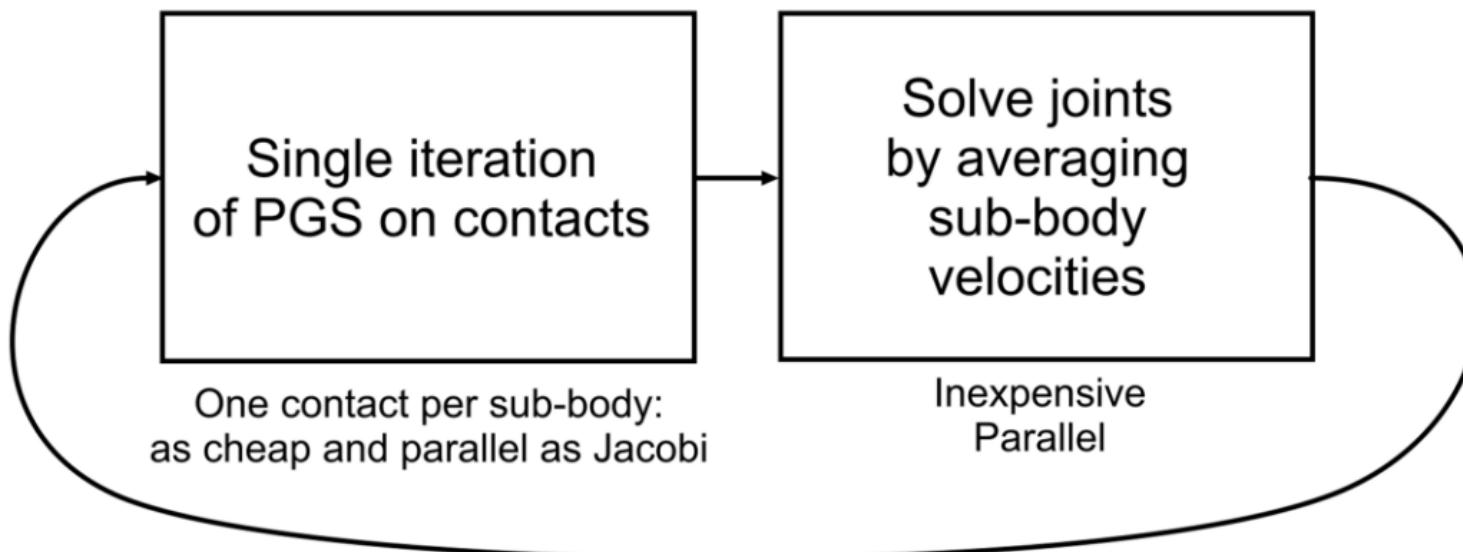


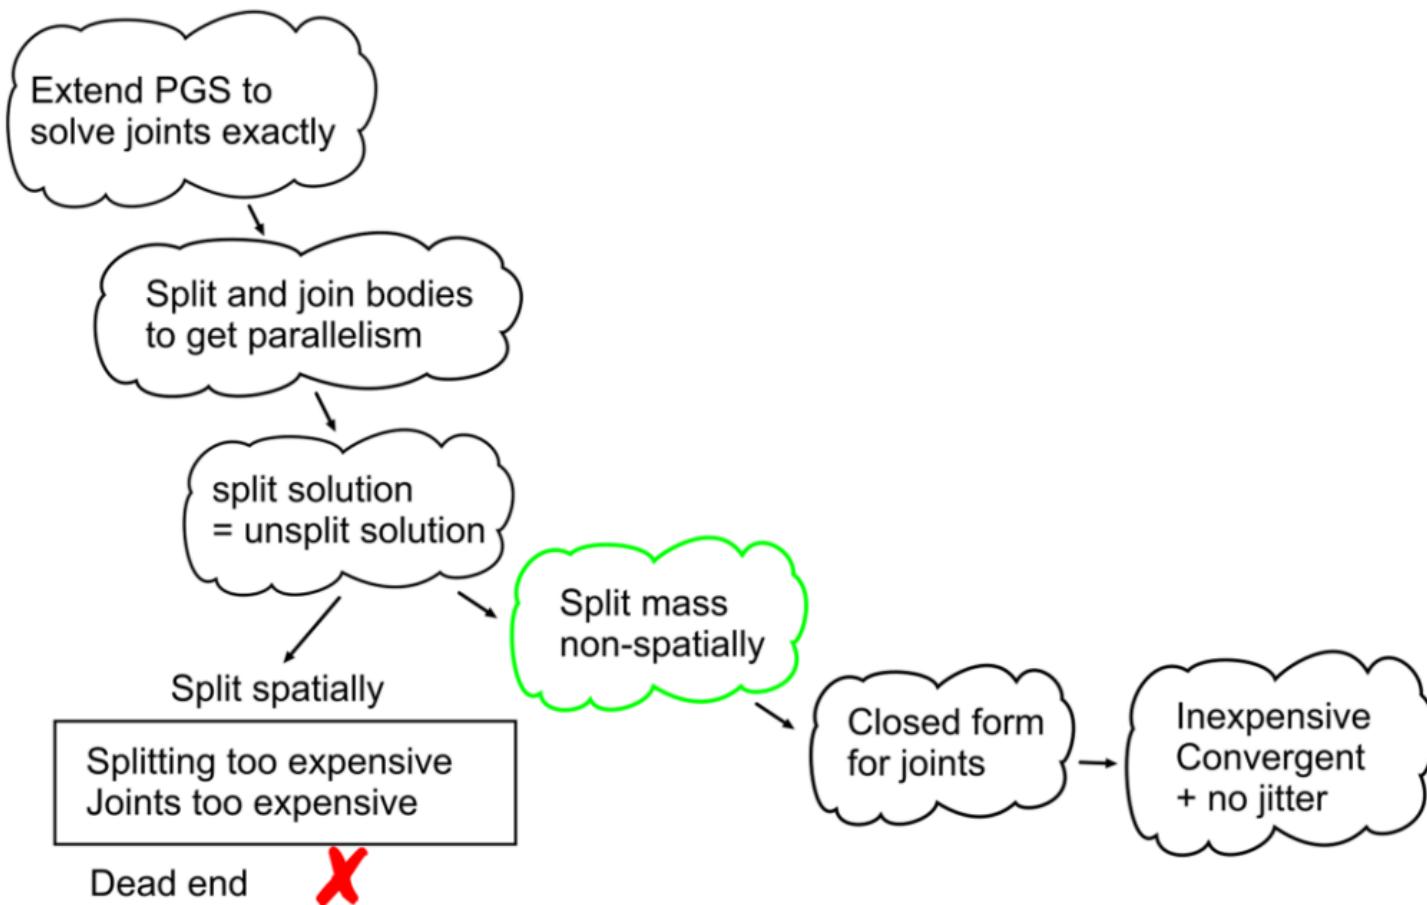
Make sub-bodies have the same spatial extent



+ two fixed joints at C.O.M.

PGS with exact joints





Section 4

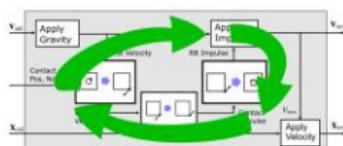
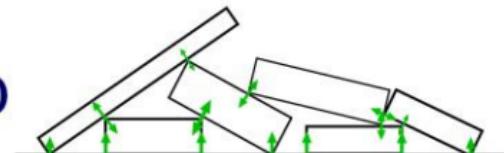
Results

Mass splitting



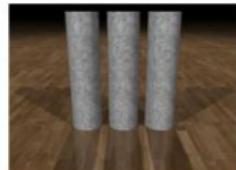
Summary

What we need rigid bodies to do

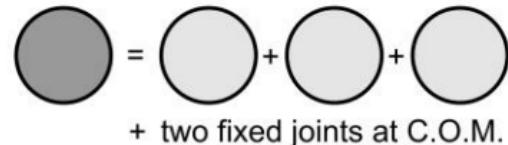


Single core solver

Parallel solver



Jitter-free



Acknowledgments

GPU rigid body technology

Feodor Benevolenski

Andrey Voroshilov

Richard Tonge

Fracture technology and demo

Matthias Müller-Fischer

Nuttapong Chentanez

Tae-Yong Kim

Aron Zoellner

Thanks also to the PhysX SDK team

NVIDIA and NVIDIA PhysX are trademarks and/or registered trademarks of NVIDIA Corporation in the United States and other countries.

